



# Designing Applications for Cloud Foundry

Writing Applications that Scale

Getting it right *By Design*

# Overview

- After completing this lesson, you should be able to:
  - Consider design considerations for Cloud Foundry

# Roadmap

- 12-Factor Applications
- Design Guidelines

# Developer Topics / Application Architecture

- Applications may require adjustments to run successfully in Cloud environment
  - See Developer Topics:
  - <http://docs.cloudfoundry.org/devguide/deploy-apps/prepare-to-deploy.html>

# Roadmap

- **12-Factor Applications**
- Design Guidelines

# 12-Factor Application

- <http://12factor.net>
- Outlines architectural principles for modern apps
  - Focus on scaling, continuous delivery, portable, and cloud ready

# 12-Factor Application

## I. Codebase

One codebase tracked in SCM, many deploys

## II. Dependencies

Explicitly declare and isolate dependencies

## III. Configuration

Store config in the environment

## IV. Backing Services

Treat backing services as attached resources

## V. Build, Release, Run

Strictly separate build and run stages

## VI. Processes

Execute app as stateless processes

## VII. Port binding

Export services via port binding

## VIII. Concurrency

Scale out via the process model

## IX. Disposability

Maximize robustness with fast startup and graceful shutdown

## X. Dev/prod parity

Keep dev, staging, prod as similar as possible

## XI. Logs

Treat logs as event streams

## XII. Admin processes

Run admin / mgmt tasks as one-off processes

# 12-Factor Application

## I. Codebase

One codebase tracked in SCM, many deploys

## II. Dependencies

Explicitly declare and isolate dependencies

## III. Configuration

Store config in the environment

- Codebase
  - An application has a single codebase
    - Multiple codebase = distributed system (not an app)
      - Each component in a codebase can (should) be an app
  - Tracked in version control
    - Git, Subversion, Mercurial, etc.
  - Multiple Deployments
    - i.e. development, testing, staging, production, etc.



# 12-Factor Application

## I. Codebase

One codebase tracked in SCM, many deploys

## II. Dependencies

Explicitly declare and isolate dependencies

## III. Configuration

Store config in the environment

- Dependencies
  - Packaged as jars (Java), RubyGems, CPAN (Perl)
  - Declared in a Manifest
    - Maven POM, Gemfile / bundle exec, etc.
  - No reliance on specific system tools
    - i.e. Linux tool not available on Windows

# 12-Factor Application

## I. Codebase

One codebase tracked in SCM, many deploys

## II. Dependencies

Explicitly declare and isolate dependencies

## III. Configuration

Store config in the environment

- Configuration
  - Separate from the code
  - Also separate from the application
    - i.e. DB credentials, hostnames, passwords
    - Acid Test – could the codebase be made open source?
    - Internal wiring (i.e. Spring configuration) considered part of codebase.
  - Environment Variables recommended.

# 12-Factor Application

## IV. Backing Services

Treat backing services as attached resources

## V. Build, Release, Run

Strictly separate build and run stages

## VI. Processes

Execute app as stateless processes

- Backing Services
  - Service consumed by app as part of normal operations
    - DB, Message Queues, SMTP servers
    - May be locally managed or third-party managed
  - Services should be treated as resources
    - Connected to via URL / configuration
    - Swappable (change in-memory DB for MySQL)

# 12-Factor Application

## IV. Backing Services

Treat backing services as attached resources

## V. Build, Release, Run

Strictly separate build and run stages

## VI. Processes

Execute app as stateless processes

- Build, Release, Run
  - Build stage – converts codebase into build (version)
    - Including managed dependencies
  - Release stage – build + config = release
    - Ready to run
  - Run – Runs app in execution environment

# 12-Factor Application

## IV. Backing Services

Treat backing services as attached resources

## V. Build, Release, Run

Strictly separate build and run stages

## VI. Processes

Execute app as stateless processes

- Processes
  - One or more discrete running processes
  - Stateless
    - Processes should not store internal state (HTTP Sessions)
  - Shared Nothing
    - Data needing to be shared should be persisted
  - Memory / local tmp storage considered volatile
  - Processes may intercommunicate via messaging / persistent storage

# 12-Factor Application

## VII. Port binding

Export services via port binding

## VIII. Concurrency

Scale out via the process model

## IX. Disposability

Maximize robustness with fast startup and graceful shutdown

- Port Binding
  - App should not need a “container”
    - Java App Server, Apache HTTPD for PHP ...
    - PaaS now takes that role
  - Apps should export HTTP as a service
    - Define as a dependency (#2)
      - Tornado (Python), Thin (Ruby), embedded Jetty/Tomcat (Java)
    - Execute at runtime
  - One App can become another App's service (#4, #6)

# 12-Factor Application

## VII. Port binding

Export services via port binding

## VIII. Concurrency

Scale out via the process model

## IX. Disposability

Maximize robustness with fast startup and graceful shutdown

- Concurrency
  - Processes are first class citizens
    - Like Unix service daemons
    - Unlike Java threads
  - Individual processes are free to multithread
    - BUT a VM can only get so large (vertical scaling).
    - Must be able to span multiple machines (horizontal scaling)

# 12-Factor Application

## VII. Port binding

Export services via port binding

## VIII. Concurrency

Scale out via the process model

## IX. Disposability

Maximize robustness with fast startup and graceful shutdown

- Disposability
  - Processes should be disposable
    - Remember, they're stateless!
  - Should be quick to start and stop
    - Should exit gracefully / finish current requests.
    - Or should be idempotent / reentrant
  - Enhances scalability and fault tolerance
  - Design *crash-only* software



# 12-Factor Application

## X. Dev/prod parity

Keep dev, staging, prod as similar as possible

## XI. Logs

Treat logs as event streams

## XII. Admin processes

Run admin / mgmt tasks as one-off processes

- Development, Staging, Production should be similar
  - Dev / Prod environments often different
    - Tool gap – devs use SQLite/Nginx, prod uses Apache/Oracle
    - Personnel gap – developers develop, admins deploy
    - Time gap - (development over weeks / months)
  - Keep differences minor
    - Reduce tool gap – use same software
    - Reduce time gap - small changes & continuous deployment
    - Reduce personnel gap - involve developers in deployment and monitoring

# 12-Factor Application

## X. Dev/prod parity

Keep dev, staging, prod as similar as possible

## XI. Logs

Treat logs as event streams

## XII. Admin processes

Run admin / mgmt tasks as one-off processes

- Logs are streams of aggregated, time-ordered events
  - Apps are not concerned with log management
    - Just write to sysout.
  - Separate log managers handle management
    - Logging as a service
- Can be managed via tools like Papertrail, Splunk ...
  - Log indexing and analysis

# 12-Factor Application

## X. Dev/prod parity

Keep dev, staging, prod as similar as possible

## XI. Logs

Treat logs as event streams

## XII. Admin processes

Run admin / mgmt tasks as one-off processes

- Admin Processes / Management Tasks Run as One-Off Processes.
  - DB Migrations, one time scripts, etc.
  - Use same environment, tools, language as application processes
    - REPL



*Read-Eval-Print Loop* = command-shell for running non-interactive shell scripts

# 12-Factor Application

## I. Codebase

One codebase tracked in SCM, many deploys

## II. Dependencies

Explicitly declare and isolate dependencies

## III. Configuration

Store config in the environment

## IV. Backing Services

Treat backing services as attached resources

## V. Build, Release, Run

Strictly separate build and run stages

## VI. Processes

Execute app as stateless processes

## VII. Port binding

Export services via port binding

## VIII. Concurrency

Scale out via the process model

## IX. Disposability

Maximize robustness with fast startup and graceful shutdown

## X. Dev/prod parity

Keep dev, staging, prod as similar as possible

## XI. Logs

Treat logs as event streams

## XII. Admin processes

Run admin / mgmt tasks as one-off processes

# Roadmap

- 12-Factor Applications
- **Design Guidelines**

# Application Architecture

- Application architecture concerns:
  - Load Balancing / Session Management
  - Local file system
  - Port Limitations


# Load Balancing Router

- CF *Router* provides automatic load balancing
  - When  $> 1$  instance
- Sticky Sessions - based on **JSESSIONID** parameter
  - Works automatically for Java Web apps
  - Other technologies need extra steps.



# Session Management

- Based on sticky sessions, managed by Router
- Session is NOT persisted between instances.
  - If an instance fails, those sessions are lost.



- ① Requests *stick* to previous session
- ② Server dies
- ③ New container & server started
- ④ New session – old session lost



# Session Management

- Session use best avoided
  - In order to achieve massive scaling
  - Easy for RESTful servers
- If Sessions are essential
  - Add persistent session management
    - For example: Gemfire cache
  - Move session-data to a light-weight persistent store
    - Such as Redis key-value store

# Local File Access

- Apps should not attempt to access the local file system
  - Short lived, not shared
- Instead, use Service abstraction when flat files are needed
  - Amazon S3, Google Cloud Storage, Dropbox, or Box
    - Examples: file-uploading
  - File Storage as a Service is coming
- Or consider using a database
  - Redis: Persistent, in-memory data
  - Mongo DB: JSON document storage

# Logging


- Loggregator will automatically handle all output logged to `sysout` or `syserr`
  - **Note:** `cf logs` receives data on port 4443 (typically blocked by corporate firewalls)
- Don't use log-files
  - Local file system is generally not available
  - Loggregator will NOT handle log files made to the file system or other sources
  - Write to `sysout` instead
  - Or consider writing log records to a fast, NoSql database
    - Can now be queried

# Resources

- All needed resources should be available via classpath
  - Example: use `classpath:` resource in Spring
- File resources not available
  - short lived / not shared
- Place configuration in `classpath:` resources
  - Spring MVC supports static web-resource in jars
    - Such as CSS, HTML, images, ...

# Port Limitations

- Port usage currently limited to SSH, HTTP and HTTPS
  - Only 22, 80, 443 open to *incoming* traffic
    - Outgoing traffic can be controlled by **Security Groups**
    - Open 4443 *inbound* in *your* firewall for logging
  - Cloud Foundry *Router* only supports these protocols



\* 4443: secure websocket for PCF logging

# Summary

- After completing this lesson, you should have learned:
  - Architectural design factors for building scalable applications in Cloud Foundry